



北京邮电大学

BEIJING UNIVERSITY OF POSTS AND TELECOMMUNICATIONS

Chapter 2: Lexical Analysis

Yanhui Guo


yhguo@bupt.edu.cn

Outline

- **The Role of Lexical Analyzer**
- Specification of Tokens (Regular Expressions)
- Recognition of Tokens (Transition Diagrams)
- The Lexical-Analyzer Generator
- Finite Automata

The Role of Lexical Analyzer

- Read the input characters of the source program, group them into lexemes, and produces a sequence of tokens
- Add lexemes into the symbol table when necessary
- Strip out comments and whitespace (blank, newline, tab etc.)
- Associate error messages with a line number (tracking newlines)



Why Separated Lexical Analysis?

- **Simplicity of compiler design**. The lexical analyzer can perform simple tasks such as dealing with comments and whitespace
- **Improved compiler efficiency**. Lexical analysis is much simpler comparing to syntax analysis
- **Higher portability of compilers**. Only lexical analyzer needs to deal with input-device-specific peculiarities (e.g., line separator)

Tokens, Patterns, and Lexemes

- A *token* is a pair $\langle \text{token name}, \text{attribute value} \rangle$
 - *Token name* is an abstract symbol representing a kind of *lexical unit*
 - *Attribute value* (optional) points to the symbol table
- A *pattern* is a description of the form that the lexemes of a token may take.
- A *lexeme* is a sequence of characters in the source program that matches the pattern for a token.
 - It is identified by the lexical analyzer as *an instance of the token*.

Examples


TOKEN	INFORMAL DESCRIPTION	SAMPLE LEXEMES
if	characters i, f	if
else	characters e, l, s, e	else
comparison	< or > or <= or >= or == or !=	<=, !=
id	letter followed by letters and digits	pi, score, D2
number	any numeric constant	3.14159, 0, 6.02e23
literal	anything but ", surrounded by "'s	"core dumped"

Consider the C statement: `printf("Total = %d\n", score);`

Lexeme	printf	score	"Total = %d\n"	(...
Token	id	id	literal	left_parenthesis	...

Attributes for Tokens

- When more than one lexeme match a pattern, the lexical analyzer must provide additional information, named *attribute values*, to the subsequent compiler phases
 - **Token names** influence parsing decisions
 - **Attribute values** influence semantic analysis, code generation etc.
- For example, an **id** token is often associated with: (1) its lexeme, (2) type, and (3) the location at which it is first found. Token attributes are stored in the *symbol table*.

A = B * 2  **<id, pointer to symbol-table entry for A>**
<assign_op>
<id, pointer to symbol-table entry for B> **<mult_op>**
<number, integer value 2>

Lexical Errors

- The lexical analyzer is unable to proceed: none of the patterns for tokens match any prefix of the remaining input
- Example: `int 3a = a * 3;`

Outline

- The Role of Lexical Analyzer
- Specification of Tokens (Regular Expressions)
- Recognition of Tokens (Transition Diagrams)
- The Lexical-Analyzer Generator
- Finite Automata

Specification of Tokens

- **Regular expression** (正则表达式, **regex** for short) is an important notation for specifying lexeme patterns
- Content of this part
 - Strings and Languages (串和语言)
 - Operations on Languages (语言上的运算)
 - Regular Expressions
 - Regular Definitions (正则定义)
 - Extensions of Regular Expressions

Strings and Languages

- **Alphabet (字母表)**: any finite set of symbols
 - Examples of symbols: **letters**, **digits**, and **punctuations**
 - Examples of alphabets: $\{1, 0\}$, **ASCII**, **Unicode**
- A **string (串)** over an alphabet is a finite sequence of symbols drawn from the alphabet
 - The length of a string s , denoted $|s|$, is the number of occurrences of symbols in s (i.e., cardinality)
- **Empty string (空串)**: the string of length 0, ϵ

Strings and Languages Cont.

- String-related terms (using **banana** for illustration)
 - **Prefix (前綴)** of string s : any string obtained by removing 0 or more symbols from the end of s (**ban**, **banana**, ϵ)
 - **Proper prefix (真前綴)**: a prefix that is not ϵ and not equal to s itself (**ban**)
 - **Suffix (后綴)**: any string obtained by removing 0 or more symbols from the beginning of s (**nana**, **banana**, ϵ).
 - **Proper suffix (真后綴)**: a suffix that is not ϵ and not equal to s itself (**nana**)

Strings and Languages Cont.

- String-related terms (using **banana** for illustration)
 - **Substring** (子串) of s : any string obtained by removing any prefix and any suffix from s (**banana**, **nan**, ϵ)
 - **Proper substring** (真子串): a substring that is not ϵ and not equal to s itself (**nan**)
 - **Subsequence** (子序列): any string formed by removing 0 or more not necessarily consecutive symbols from s



How many substrings does **banana** have?

Strings and Languages Cont.

- String-related operations (串的运算)
 - **Concatenation (连接)**: the concatenation of two strings x and y , denoted xy , is the string formed by appending y to x
 - $x = \text{dog}, y = \text{house}, xy = \text{doghouse}$
 - **Exponentiation (幂/指数运算)**: $s^0 = \epsilon, s^1 = s, s^i = s^{i-1}s$
 - $x = \text{dog}, x^0 = \epsilon, x^1 = \text{dog}, x^3 = \text{dogdogdog}$

Strings and Languages Cont.

- A **language** (语言) is any **countable set**¹ of strings over some fixed alphabet
 - The set containing only the empty string, that is $\{\epsilon\}$, is a language, denoted \emptyset
 - The set of all **grammatically correct English sentences**
 - The set of all **syntactically well-formed C programs**

¹In mathematics, a countable set is a set with the same cardinality (number of elements) as some subset of the set of natural numbers. A countable set is either a finite set or a countably infinite set.

Strings and Languages Cont.



Stephen C. Kleene

- Operations on Languages (语言的运算)
 - 并, 连接, Kleene闭包, 正闭包

OPERATION	DEFINITION AND NOTATION
<i>Union of L and M</i>	$L \cup M = \{s \mid s \text{ is in } L \text{ or } s \text{ is in } M\}$
<i>Concatenation of L and M</i>	$LM = \{st \mid s \text{ is in } L \text{ and } t \text{ is in } M\}$
<i>Kleene closure of L</i>	$L^* = \bigcup_{i=0}^{\infty} L^i$
<i>Positive closure of L</i>	$L^+ = \bigcup_{i=1}^{\infty} L^i$

https://en.wikipedia.org/wiki/Stephen_Cole_Kleene

Strings and Languages Cont.

- Operations on Languages (Example)
 - $L = \{A, B, \dots, Z, a, b, \dots, z\}$
 - $D = \{0, 1, \dots, 9\}$
 - $L \cup D = \{A, B, \dots, Z, a, b, \dots, z, 0, 1, \dots, 9\}$
 - LD : the set of 520 strings of length two, each consisting of one letter followed by one digit
 - L^4 : the set of all 4-letter string
 - L^* : the set of all strings of letters, including ϵ
 - $L(L \cup D)^*$: ?
 - D^+ : ?

Regular Expressions

Rules that define regexps over an alphabet Σ :

- BASIS: two rules form the basis:
 - ϵ is a regexp, $L(\epsilon) = \{\epsilon\}$
 - If a is a symbol in Σ , then a is a regexp, and $L(a) = \{a\}$
- INDUCTION: Suppose r and s are regexps denoting the languages $L(r)$ and $L(s)$
 - $(r) \mid (s)$ is a regexp denoting the language $L(r) \cup L(s)$
 - $(r)(s)$ is a regexp denoting the language $L(r)L(s)$
 - $(r)^*$ is a regexp denoting $(L(r))^*$
 - (r) is a regexp denoting $L(r)$. Additional parentheses do not change the language an expression denotes.

Regular Expressions Cont.

- Following the rules, regexps often contain **unnecessary pairs of parentheses**. We may drop some if we adopt the conventions:
 - **Precedence**: **closure**^{*} > **concatenation** > **union** |
 - **Associativity**: All three operators are **left associative**, meaning that operations are grouped from the left, e.g.,
 $a \mid b \mid c$ would be interpreted as $(a \mid b) \mid c$
- Example: $(a) \mid ((b)^*(c)) = a \mid b^*c$

Regular Expressions Cont.

- Examples: Let $\Sigma = \{a, b\}$
 - $a|b$ denotes the language $\{a, b\}$
 - $(a|b)(a|b)$ denotes $\{aa, ab, ba, bb\}$
 - a^* denotes $\{\epsilon, a, aa, aaa, \dots\}$
 - $(a|b)^*$ denotes the set of all strings consisting of 0 or more a's or b's: $\{\epsilon, a, b, aa, ab, ba, bb, aaa, \dots\}$
 - $a|a^*b$ denotes the string a and all strings consisting of 0 or more a's and ending in b: $\{a, b, ab, aab, aaab, \dots\}$

Regular Expressions Cont.

- A **regular language** (正则语言) is a language that can be defined by a regexp
- If two regexps r and s denote the same language, they are **equivalent**, written as $r = s$
- Each **algebraic law** below asserts that expressions of two different forms are equivalent

LAW	DESCRIPTION
$r s = s r$	$ $ is commutative
$r (s t) = (r s) t$	$ $ is associative
$r(st) = (rs)t$	Concatenation is associative
$r(s t) = rs rt; (s t)r = sr tr$	Concatenation distributes over $ $
$\epsilon r = r\epsilon = r$	ϵ is the identity for concatenation
$r^* = (r \epsilon)^*$	ϵ is guaranteed in a closure
$r^{**} = r^*$	$*$ is idempotent

$(a|b)(a|b)$
=
 $aa|ab|ba|bb$
Is it true???

Regular Definitions

- For **notational convenience**, we can give names to certain regexps and use those names in subsequent expressions
- If Σ is an alphabet of basic symbols, then a **regular definition** is a sequence of definitions of the form:

$$d_1 \rightarrow r_1$$

$$d_2 \rightarrow r_2$$

...

$$d_n \rightarrow r_n$$

where:

- Each d_i is a new symbol not in Σ and not the same as the other d 's
- Each r_i is a regexp over the alphabet $\Sigma \cup \{d_1, d_2, \dots, d_{i-1}\}$

Regular Definitions - Examples

- **Regex** for **C** identifiers

letter_ → A | B | ... | Z | a | b | ... | z | _
digit → 0 | 1 | ... | 9
id → *letter_* (*letter_* | *digit*)*

_hello valid?
3times valid?

- **Regex** for **C** identifiers

$(A|B|\dots|Z|a|b|\dots|z|_)((A|B|\dots|Z|a|b|\dots|z|_)|(0|1|\dots|9))^*$

Extensions of Regular Expressions

- **Basic operators:** union $|$, concatenation, and Kleene closure^{*} (proposed by Kleene in 1950s)
- A few **notational extensions**:
 - **One of more instances:** the unary, postfix operator $+$
 - $r^+ = rr^*$, $r^* = r^+ | \epsilon$
 - **Zero or one instance:** the unary postfix operator $?$
 - $r? = r | \epsilon$
 - **Character classes:** shorthand for a logical sequence
 - $[a_1a_2\dots a_n] = a_1 | a_2 | \dots | a_n$
 - $[a-e] = a | b | c | d | e$
- The extensions are **only for notational convenience**, they do not change the descriptive power of regexps

Outline

- The Role of Lexical Analyzer
- Specification of Tokens (Regular Expressions)
- Recognition of Tokens (Transition Diagrams)
- The Lexical-Analyzer Generator
- Finite Automata

Recognition of Tokens

- Lexical analyzer examines the input string and finds a prefix that matches one of the tokens
- The first thing when building a lexical analyzer is to define the patterns of tokens using regular definitions
- **A special token: ws** \rightarrow (blank | tab | newline)⁺
 - When the lexical analyzer recognizes the **whitespace token**, it does not return it to the parser, but rather restart the lexical analysis from the character that follows the whitespace (the next token gets returned to the parser)

Example: Patterns and Tokens

$digit \rightarrow [0-9]$
 $digits \rightarrow digit^+$
 $number \rightarrow digits (. digits)? (E [+ -]? digits)?$
 $letter \rightarrow [A-Za-z]$
 $id \rightarrow letter (letter | digit)^*$
 $if \rightarrow if$
 $then \rightarrow then$
 $else \rightarrow else$
 $relop \rightarrow < | > | <= | >= | = | <>$


Patterns for tokens

LEXEMES	TOKEN NAME	ATTRIBUTE VALUE
Any <i>ws</i>	–	–
if	if	–
then	then	–
else	else	–
Any <i>id</i>	id	Pointer to table entry
Any <i>number</i>	number	Pointer to table entry
<	relop	LT
<=	relop	LE
=	relop	EQ
<>	relop	NE
>	relop	GT
>=	relop	GE

Tokens, their patterns, and attribute values

Transition Diagrams (状态转换图)

- An important step in constructing a lexical analyzer is to convert patterns into “**transition diagrams**”
- Transition diagrams have a collection of nodes, called **states** (状态) and **edges** (边) directed from one node to another




LEXEMES	TOKEN NAME	ATTRIBUTE VALUE
<	relop	LT
<=	relop	LE
=	relop	EQ
<>	relop	NE
>	relop	GT
>=	relop	GE

The transition diagram in the left is for recognizing **relop** tokens

States

- Represent conditions that could occur during the process of scanning the input for recognizing lexemes of tokens (what characters we have seen)
- The **start state** (开始状态), or initial state, is indicated by an edge labeled “start”, which enters from nowhere
- Certain states are said to be **accepting** (接受状态), or final, indicating that a lexeme has been found





States 2, 3, 4, 5, 7, 8 are accepting.

By convention, we indicate accepting states by **double circles**

States Cont.



- At certain accepting states, the found lexeme may not contain all characters that we have seen from the start state
- Such states are annotated with *
- When entering * states, it is necessary to **retract** (回退) the forward pointer, which points to the next char in the input string)



We should retract forward one position in the above case

Edges


- **Edges** are directed from one state to another
- Each edge is labeled by a symbol or set of symbols



In the above case, we should follow the circled edge to enter state 3 and advance the forward pointer


Recognition of Reserved Words and Identifiers (保留字和标识符的识别)

- In many languages, **reserved** words or **keywords** (e.g., then) also match the pattern of identifiers
- **Problem**: the transition diagram that searches for identifiers can also recognize reserved words




Handling Reserved Words

- **Strategy 1:** Preinstall the reserved words in the symbol table. Put a field in the symbol-table entries to indicate that these strings are not ordinary identifiers



- **Strategy 2:** Create a separate transition diagram with a high priority for each keyword




Building a Lexical Analyzer from Transition Diagrams

```


TOKEN getRelop()
{
    TOKEN retToken = new(RELOP);
    while(1) { /* repeat character processing until a return
                or failure occurs */
        switch(state) {
            case 0: c = nextChar();
                    if ( c == '<' ) state = 1;
                    else if ( c == '=' ) state = 5;
                    else if ( c == '>' ) state = 6;
                    else fail(); /* lexeme is not a relop */
                    break;
            case 1: ...
            ...
            case 8: retract();
                    retToken.attribute = GT;
                    return(retToken);
        }
    }
}

```



Building a Lexical Analyzer from Transition Diagrams

```
TOKEN getRelop()
{
    TOKEN retToken = new(RELOP);
    while(1) { /* repeat character processing until a return
               or failure occurs */
        switch(state){
            case 0: c = nextChar();
                    if ( c == '<' ) state = 1;
                    else if ( c == '=' ) state = 5;
                    else if ( c == '>' ) state = 6;
                    else fail(); /* lexeme is not a relop */
                    break;
            case 1: ...
            ...
            case 8: retract();
                    retToken.attribute = GT;
                    return(retToken);
        }
    }
}
```




Use a variable state to record the current state

Building a Lexical Analyzer from Transition Diagrams

```

TOKEN getRelop()
{
    TOKEN retToken = new(RELOP);
    while(1) { /* repeat character processing until a return
               or failure occurs */
        switch(state) {
            case 0: c = nextChar();
                    if ( c == '<' ) state = 1;
                    else if ( c == '=' ) state = 5;
                    else if ( c == '>' ) state = 6;
                    else fail(); /* lexeme is not a relop */
                    break;
            case 1: ...
            ...
            case 8: retract();
                    retToken.attribute = GT;
                    return(retToken);
        }
    }
}
    
```




A switch statement based on the value of state takes us to code for each of the possible states

Sketch of implementation of relop transition diagram

Building a Lexical Analyzer from Transition Diagrams

```
TOKEN getRelop()
{
    TOKEN retToken = new(RELOP);
    while(1) { /* repeat character processing until a return
                or failure occurs */
        switch(state) {
            case 0: c = nextChar();
                    if ( c == '<' ) state = 1;
                    else if ( c == '=' ) state = 5;
                    else if ( c == '>' ) state = 6;
                    else fail(); /* lexeme is not a relop */
                    break;
            case 1: ...
            ...
            case 8: retract();
                    retToken.attribute = GT;
                    return(retToken);
        }
    }
}
```



The code of a normal state:

1. Read the next character
2. Determine the next state
3. If step 2 fails, do error recovery

Sketch of implementation of relop transition diagram


Building a Lexical Analyzer from Transition Diagrams

```

TOKEN getRelop()
{
    TOKEN retToken = new(RELOP);
    while(1) { /* repeat character processing until a return
                or failure occurs */
        switch(state) {
            case 0: c = nextChar();
                    if ( c == '<' ) state = 1;
                    else if ( c == '=' ) state = 5;
                    else if ( c == '>' ) state = 6;
                    else fail(); /* lexeme is not a relop */
                    break;

            case 1: ...

            case 8: retract();
                   retToken.attribute = GT;
                   return(retToken);
        }
    }
}
    
```



The code of an accepting state:

1. Perform retraction if the state has *
2. Set token attribute values
3. Return the token to parser

Sketch of implementation of relop transition diagram

Building the Entire Lexical Analyzer

- **Strategy 1:** Try the transition diagram for each token sequentially
 - fail() resets the pointer forward and starts the next diagram
- **Strategy 2:** Run transition diagrams in parallel
 - Need to resolve the case where one diagram finds a lexeme and others are still able to process input.
 - **Solution:** take the longest prefix of the input that matches any pattern
- **Strategy 3: Combining all transition diagrams into one (preferred)**
 - Allow the transition diagram to read input until there is no possible next state
 - Take the longest lexeme that matched any pattern

Outline

- The Role of Lexical Analyzer
- Specification of Tokens (Regular Expressions)
- Recognition of Tokens (Transition Diagrams)
- The Lexical-Analyzer Generator
- Finite Automata

The Lexical-Analyzer Generator Lex

- Lex, or a more recent tool Flex, allows one to specify a lexical analyzer by specifying regexps to describe patterns for tokens
- Often used with Yacc/Bison to create the frontend of compiler

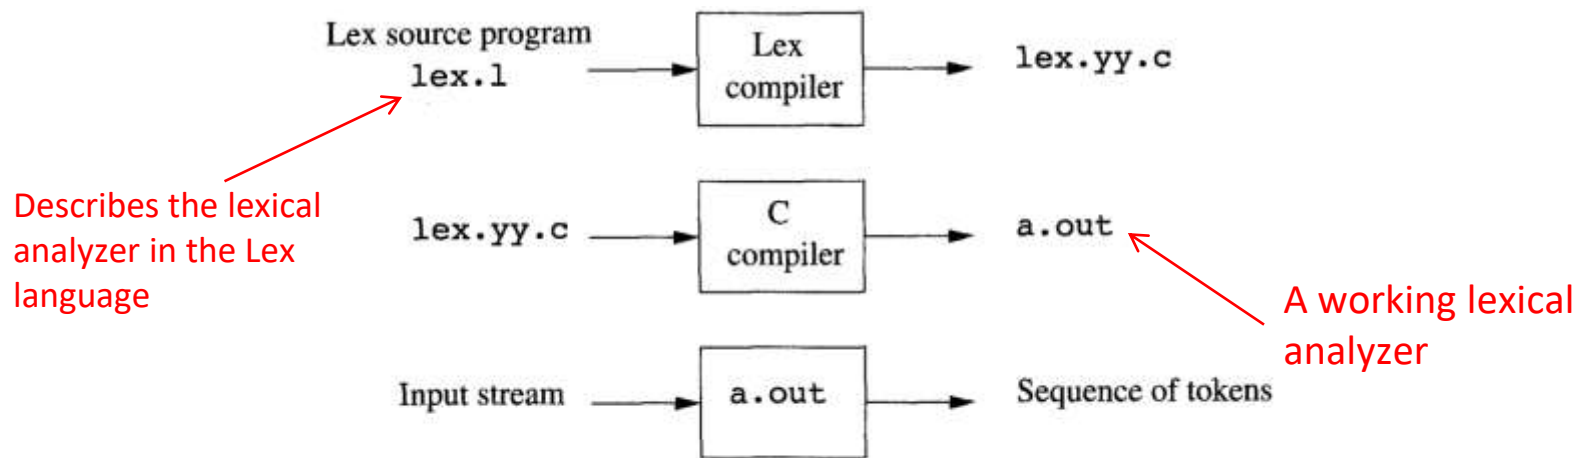


Figure 3.22: Creating a lexical analyzer with Lex

Structure of Lex Programs

- A Lex program has three sections separated by %%
 - **Declaration (声明)**
 - Variables, manifest constants (e.g., token names)
 - Regular definitions
 - **Translation rules (转换规则)** in the form “Pattern {Action}”
 - Each **pattern (模式)** is a regexp (may use the regular definitions of the declaration section)
 - **Actions (动作)** are fragments of code, typically in C, which are executed when the pattern is matched
 - **Auxiliary functions section (辅助函数)**
 - Additional functions that can be used in the actions

Lex Program Example

```
%{  
    /* definitions of manifest constants  
    LT, LE, EQ, NE, GT, GE,  
    IF, THEN, ELSE, ID, NUMBER, RELOP */  
%}  
  
/* regular definitions */  
delim      [ \t\n]  
ws         {delim}+  
letter     [A-Za-z]  
digit      [0-9]  
id         {letter}({letter}|{digit})*  
number     {digit}+(\.{digit}+)?(E[+-]?{digit}+)?  
  
%%
```

Anything in between %{ and}% is copied directly to lex.yy.c.

In the example, there is only a comment, not real C code to define manifest constants

Regular definitions that can be used in translation rules

Section separator

Lex Program Example Cont.

```
{ws}      { /* no action and no return */ }
if        { return(IF); }
then      { return(THEN); }
else      { return(ELSE); }
{id}      { yylval = (int) installID(); return(ID); }
{number}  { yylval = (int) installNum(); return(NUMBER); }
"<"       { yylval = LT; return(RELOP); }
"<="      { yylval = LE; return(RELOP); }
"="       { yylval = EQ; return(RELOP); }
"<>"      { yylval = NE; return(RELOP); }
">"       { yylval = GT; return(RELOP); }
">="      { yylval = GE; return(RELOP); }

%%
```

Continue to recognize other tokens

Return token name to the parser

Place the lexeme found in the symbol table

A global variable that stores a pointer to the symbol table entry for the lexeme. Can be used by the parser or a later component of the compiler.

Lex Program Example Cont.

- Everything in the auxiliary function section is copied directly to the file **lex.yy.c**
- Auxiliary functions may be used in actions in the translation rules

```
int installID() { /* function to install the lexeme, whose
                  first character is pointed to by yytext,
                  and whose length is yyleng, into the
                  symbol table and return a pointer
                  thereto */
}
```

Variables defined and set automatically by the lexical analyzer Lex generates

```
int installNum() { /* similar to installID, but puts numerical
                    constants into a separate table */
}
```

Conflict Resolution

- When the generated lexical analyzer runs, it analyzes the input looking for **prefixes that match any of its patterns**.^{*}
- **Rule 1:** If it finds multiple such prefixes, it takes the **longest** one
 - The analyzer will treat `<=` as a single lexeme, rather than `<` as one lexeme and `=` as the next
- **Rule 2:** If it finds a prefix matching different patterns, **the pattern listed first** in the Lex program is chosen.
 - If the keyword patterns are listed before identifier pattern, the lexical analyzer will not recognize keywords as identifiers

^{*}See Flex manual for details (Chapter 8: How the input is matched) at <http://dinosaur.compilertools.net/flex/>

Outline

- The Role of Lexical Analyzer
- Specification of Tokens (Regular Expressions)
- Recognition of Tokens (Transition Diagrams)
- The Lexical-Analyzer Generator
- Finite Automata →
 - NFA & DFA
 - NFA → DFA
 - Regexp → NFA
 - Combining NFA's (to be discussed in lab)
 - DFA Minimization (to be discussed in lab)

Finite Automata (有穷自动机)


- Finite automata are the simplest machines to recognize patterns
- **They are essentially graphs** like transition diagrams. They simply say “yes” or “no” about each possible input string.
 - **Nondeterministic finite automata (NFA, 非确定有穷自动机)**: A symbol can label several edges out of the same state (allowing multiple target states), and the empty string ϵ is a possible label.
 - **Deterministic finite automata (DFA, 确定有穷自动机)**: For each state and for each symbol of its input alphabet, there is exactly one edge with that symbol leaving that state.
- NFA and DFA recognize the same languages, **regular languages**, that regexps can describe.

Nondeterministic Finite Automata


- A nondeterministic finite automaton (NFA) consists of:
 - A finite set of states S
 - A set of input symbols Σ , the **input alphabet**. We assume that the empty string ϵ is never a member of Σ .
 - A **transition function** that gives, for each state, and for each symbol in $\Sigma \cup \{\epsilon\}$ **a set of next states**.
 - A **start state** (or initial state) s_0 from S
 - A set of **accepting states** (or final states) F , a subset of S

NFA Example

- $S = \{0, 1, 2, 3\}$
- Start state: 0
- Accepting states: $\{3\}$
- Transition function
 - $(0, a) \rightarrow \{0, 1\}$ $(0, b) \rightarrow \{0\}$
 - $(1, b) \rightarrow \{2\}$ $(2, b) \rightarrow \{3\}$



Transition Table




- Another representation of an NFA
 - **Rows** correspond to states
 - **Columns** correspond to the input symbols or ϵ
 - **The table entry** for a state-input pair lists the set of next states
 - \emptyset : the transition function has no info about the state-input pair

STATE	a	b	ϵ
0	$\{0, 1\}$	$\{0\}$	\emptyset
1	\emptyset	$\{2\}$	\emptyset
2	\emptyset	$\{3\}$	\emptyset
3	\emptyset	\emptyset	\emptyset

Acceptance of Input Strings


- An NFA **accepts** an input string **x** **if and only if**
 - There is a path in the transition graph from the start state to one accepting state, such that the symbols along the path form x (ϵ labels are ignored).




accepts the string aabb

- The **language** defined or accepted by an NFA
 - The set of strings labelling some path from the start state to an accepting state

NFA and Regular Languages



$L((a|b)^*abb)$



$L(aa^* | bb^*)$

Deterministic Finite Automata (DFA)

- A deterministic finite automaton is a special NFA where:
 - There are no moves on input ϵ
 - For each state s and input symbol a , there is exactly one edge out of s labeled a
- DFA can efficiently accept/reject strings (recognize patterns)
- Every regexp and every NFA can be converted to a DFA accepting the same language

Simulating a DFA

- **Input:**


- String x terminated by an end-of-file character eof.
- DFA D with start state s_0 , accepting states F , and transition function $move$

- **Output:** Answer “yes” if D accepts x ; “no” otherwise

```
 $s = s_0;$   
 $c = nextChar();$   
while (  $c \neq eof$  ) {  
     $s = move(s, c);$   
     $c = nextChar();$   
}  
if (  $s$  is in  $F$  ) return "yes";  
else return "no";
```

DFA Example

- Given the input string ababb, the DFA below enters the sequence of states 0, 1, 2, 1, 2, 3 and returns "yes"



What's the language defined by this DFA?

From Regular Expressions to Automata

- Regexps concisely & precisely describe the patterns of tokens
- DFA can efficiently recognize patterns (comparatively, the simulation of NFA is less straightforward*)
- When implementing lexical analyzers, regexps are often converted to DFA:
 - $\text{Regexp} \rightarrow \text{NFA} \rightarrow \text{DFA}$


* There may be multiple transitions at a state when seeing a symbol

Conversion of an NFA to a DFA

- The [subset construction](#) algorithm (子集构造法)
 - **Insight:** Each state of the constructed DFA corresponds to a set of NFA states
 - After reading the input $a_1a_2\dots a_n$, the DFA is in the state which corresponds to the set of states that the NFA can reach, from its start state, following paths labeled $a_1a_2\dots a_n$
 - The algorithm simulates “in parallel” all possible moves an NFA can make on a given input string

Example for Algorithm Illustration

- The NFA below accepts the string babb
 - There exists a path from the start state 0 to the accepting state 10, on which the labels on the edges form the string babb



Subset Construction Technique

- It is possible that # DFA states is **exponential** in # NFA states (**worst case**)
 - Each DFA state corresponds to a subset of NFA states
- However, for real languages, the NFA and DFA have **approximately the same number of states**, and the exponential behavior is not seen

Subset Construction Technique Cont.

- Operations used in the algorithm:
 - **ϵ -closure(s)**: Set of NFA states reachable from NFA state s on ϵ -transitions alone
 - **ϵ -closure(T)**: Set of NFA states reachable from some NFA state s in set T on ϵ -transitions alone
 - That is, $\bigcup_{s \in T} \epsilon\text{-closure}(s)$
 - **move(T, a)**: Set of NFA states to which there is a transition on input symbol a from some state s in T

Subset Construction Technique Cont.

- Computing ϵ -closure(T)

- It is a graph traversal process (only consider ϵ edges)
- Computing ϵ -closure(s) is essentially the same

```
push all states of  $T$  onto stack;  
initialize  $\epsilon$ -closure( $T$ ) to  $T$ ;  
while ( stack is not empty ) {  
    pop  $t$ , the top element, off stack;  
    for ( each state  $u$  with an edge from  $t$  to  $u$  labeled  $\epsilon$  )  
        if (  $u$  is not in  $\epsilon$ -closure( $T$ ) ) {  
            add  $u$  to  $\epsilon$ -closure( $T$ );  
            push  $u$  onto stack;  
        }  
}
```


Subset Construction Technique Cont.

- The construction of the set of D 's states, **Dstates**, and the transition function **Dtran** is also a search process
 - Initially, the only state in $Dstates$ is ϵ -closure(s_0) and it is unmarked
 - Unmarked state means that its next states have not been explored

```
while ( there is an unmarked state  $T$  in  $Dstates$  ) {  
    mark  $T$ ;  
    for ( each input symbol  $a$  ) {  
         $U = \epsilon$ -closure( $move(T, a)$ );  
        if (  $U$  is not in  $Dstates$  )  
            add  $U$  as an unmarked state to  $Dstates$ ;  
         $Dtran[T, a] = U$ ;  
    }  
}
```

Example

- **A**: ϵ -closure(0) = {0, 1, 2, 4, 7}
- **B**: $\text{Dtran}[A, a] = \epsilon$ -closure({3, 8}) = {1, 2, 3, 4, 6, 7, 8}
- **C**: $\text{Dtran}[A, b] = \epsilon$ -closure({5}) = {1, 2, 4, 5, 6, 7}
- **D**: $\text{Dtran}[B, b] = \epsilon$ -closure({5, 9}) = {1, 2, 4, 5, 6, 7, 9}
- ...



Transition Table of the DFA

- Start state: A; Accepting states: {E}

NFA STATE	DFA STATE	<i>a</i>	<i>b</i>
{0, 1, 2, 4, 7}	<i>A</i>	<i>B</i>	<i>C</i>
{1, 2, 3, 4, 6, 7, 8}	<i>B</i>	<i>B</i>	<i>D</i>
{1, 2, 4, 5, 6, 7}	<i>C</i>	<i>B</i>	<i>C</i>
{1, 2, 4, 5, 6, 7, 9}	<i>D</i>	<i>B</i>	<i>E</i>
{1, 2, 3, 5, 6, 7, 10}	<i>E</i>	<i>B</i>	<i>C</i>

This DFA can be further minimized: A and C have the same moves on all symbols and can be merged.


Regular Expression to NFA

- **Thompson's construction algorithm** (Thompson构造法)
- The algorithm works **recursively** by splitting an expression into its constituent subexpressions, from which the NFA will be constructed using a set of rules
- The rules for constructing an NFA:
 - **Two basis rules (基本规则)**: handle subexpressions with no operators
 - **Three inductive rules (归纳规则)**: construct larger NFA's from the NFA's for subexpressions


Thompson's Construction Algorithm

Two basis rules:

1. The **empty expression** ϵ is converted to




2. Any subexpression a (a **single symbol** in input alphabet) is converted to



Thompson's Construction Algorithm

The inductive rules: the union case


- $s \mid t$: $N(s)$ and $N(t)$ are NFA's for subexpressions s and t



Thompson's Construction Algorithm

The inductive rules: the concatenation case

- **st** : $N(s)$ and $N(t)$ are NFA's for subexpressions s and t




Merging the accepting state of $N(s)$ and the start state of $N(t)$

Thompson's Construction Algorithm

The inductive rules: the Kleene star case


- s^* : $N(s)$ is the NFA for subexpression s




Example

- Use Thompson's algorithm to construct an NFA for the regexp
 $r = (a|b)^*abb$

NFA for the first a (apply basis rule #1)




NFA for the first b (apply basis rule #1)




Example Cont.

- NFA for $(a|b)$ (apply inductive rule #1)




Example Cont.

- NFA for $(a|b)^*$ (apply inductive rule #3)




Example Cont.

- NFA for the second **a** (apply basic rule #1)




Example Cont.

- NFA for $(a|b)^*a$ (apply inductive rule #2) ...



Combining NFA's

- **Why?** In the lexical analyzer, we need a single automaton to recognize lexemes matching any pattern (in the lex program)
- **How?** Introduce a new start state with ϵ -transitions to each of the start states of the NFA's N_i for pattern p_i




- The language that can be accepted by the big NFA is the union of the languages that can be accepted by the small NFA's
- Different accepting states correspond to different patterns

DFA's for Lexical Analyzers

- Convert the NFA for all the patterns into an equivalent DFA, using the subset construction algorithm
- An accepting state of the DFA corresponds to a subset of the NFA states, in which at least one is an accepting NFA state
 - If there are more than one accepting NFA state, this means that conflicts arise (the prefix of the input string matches multiple patterns)
 - Upon conflicts, find the first pattern whose accepting state is in the set and make that pattern the output of the DFA state

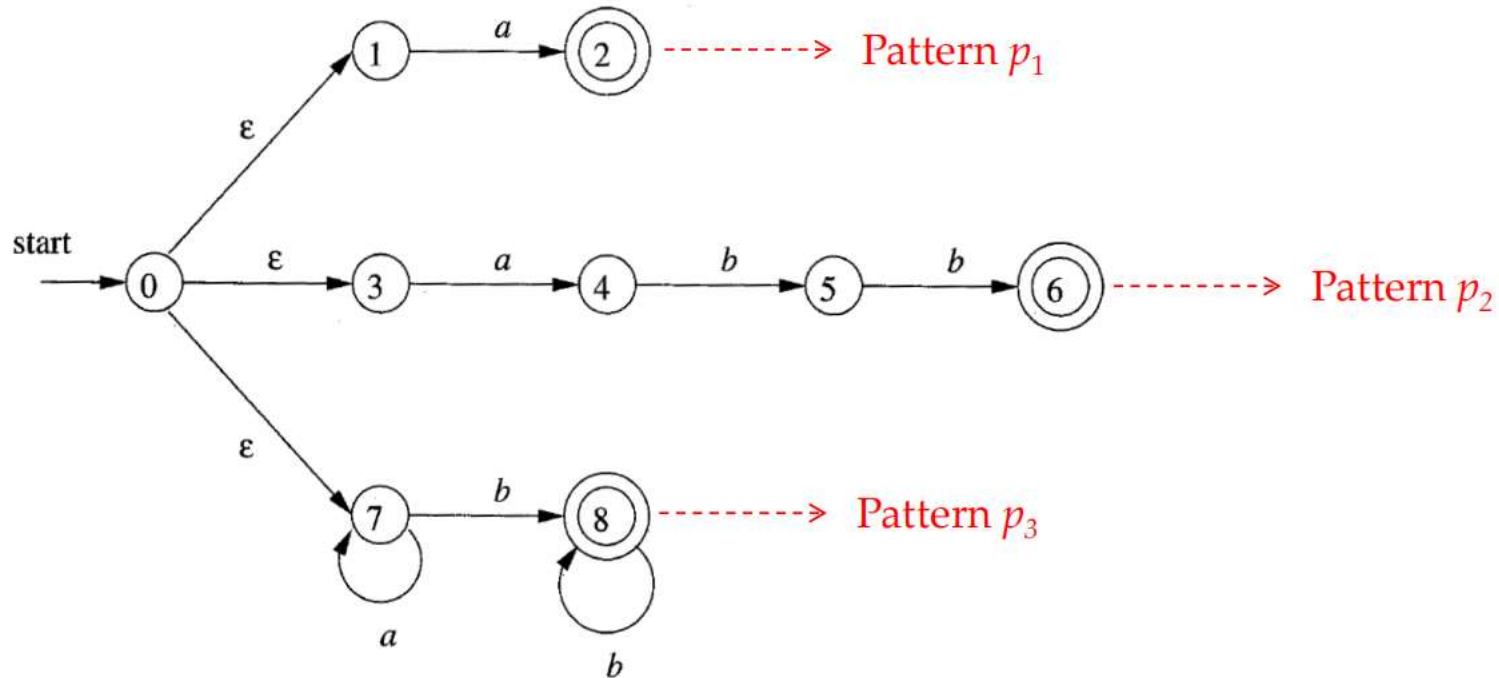
Example

- Suppose we have three patterns: p_1 , p_2 , and p_3
 - **a** {action A_1 for pattern p_1 }
 - **abb** {action A_2 for pattern p_2 }
 - **a^*b^+** {action A_3 for pattern p_3 }
- We first construct an NFA for each pattern




Example Cont.

- Combining the three NFA's



Example Cont.


- Converting the big NFA to a DFA




6 and 8 are two accepting NFA states corresponding to two patterns. We choose Pattern p_2 , which is specified before p_3

Minimizing # States of a DFA

- There can be many DFA's recognizing the same language




Two equivalent DFA's, both recognizing regular language $L((a|b)^*abb)$



Minimizing # States of a DFA


- There can be many DFA's recognizing the same language



States A and C are equivalent:

Neither is accepting, and on any symbol they transfer to the same state

A and C behave like 123



Minimizing # States of a DFA Cont.


- There is always a unique minimum-state DFA for any regular language (state name does not matter)
- The minimum-state DFA can be constructed from any DFA for the same language by grouping sets of equivalent states

Distinguishing States

- **Distinguishable states**

- We say that string x distinguishes state s from state t if **exactly one** of the states reached from s and t by following the path with label x is an accepting state
- States s and t are **distinguishable** if there exists some string that distinguishes them

- **Indistinguishable states** are equivalent and can be merged



- The empty string ϵ distinguishes any accepting state from any nonaccepting state
- The string bb distinguishes state A from B , since bb takes A to a nonaccepting state C , but takes B to accepting state E

DFA State-Minimization Algorithm

Works by partitioning the states of a DFA into groups of states that cannot be distinguished (an iterative process)

- The algorithm maintains a partition (划分), whose groups are sets of states that have not yet been distinguished
- Any two states from different groups are known to be distinguishable
- When the partition cannot be refined further by breaking any group into smaller groups, we have the minimum-state DFA

The Partitioning Part

1. Start with an initial partition Π with two groups, F and S-F, the accepting and nonaccepting states of D
2. Apply the procedure below to construct a new partition Π_{new}

```
initially, let  $\Pi_{\text{new}} = \Pi$ ;  
for ( each group  $G$  of  $\Pi$  ) {  
    partition  $G$  into subgroups such that two states  $s$  and  $t$   
        are in the same subgroup if and only if for all  
        input symbols  $a$ , states  $s$  and  $t$  have transitions on  $a$   
        to states in the same group of  $\Pi$ ;  
    /* at worst, a state will be in a subgroup by itself */  
    replace  $G$  in  $\Pi_{\text{new}}$  by the set of all subgroups formed;  
}
```

3. If $\Pi_{\text{new}} == \Pi$, let $\Pi_{\text{final}} = \Pi$ and the partitioning finishes; Otherwise, $\Pi = \Pi_{\text{new}}$ and repeat step 2


The Construction Part

Choose one state in each group of Π_{final} as the **representative** for that group. The representatives will be the states of the minimum state DFA D'

- The start state of D' is the representative of the group containing the start state of D
- The accepting states of D' are the representatives of those groups that contain an accepting state of D
- Establish transitions:
 - Let s be the representative of group G in Π_{final} ; Let the transition of D from s on input a be to state t ; Let r be the representative of t 's group H
 - Then in D' , there is a transition from s to r on input a

Example

- Initial partition: $\{A, B, C, D\}, \{E\}$
- Handling group $\{A, B, C, D\}$: b splits it to two subgroups $\{A, B, C\}$ and $\{D\}$
- Handling group $\{A, B, C\}$: b splits it to two subgroups $\{A, C\}$ and $\{B\}$
- Picking A, B, D, E as representatives to construct the minimum-state DFA



STATE	a	b
A	B	A
B	B	D
D	B	E
E	B	A

State Minimization in Lexical Analyzers

The basic idea is the same as the state-minimization algorithm for DFA.

- Differences are:
 - Each accepting state in the lexical analyzer's DFA corresponds to a different pattern. These states are not equivalent.
 - So the initial partition should be: one group of non accepting states + groups of accepting states for different patterns

Example

Initial partition: $\{0137, 7\}, \{247\}, \{68\}, \{8, 58\}, \{\emptyset\}$

- We add a dead state \emptyset : we suppose has transitions to itself on inputs a and b. It is also the target of missing transitions on a from states 8, 58, and 68.

